

Searching in a sorted array

Algorithms and data structures ID1021

Johan Montelius

Fall term 2023

Introduction

Searching for a key in an unsorted array is as you probably now know quite expensive. If the elements in the array are not sorted, the only way to find an element is to go through the whole data structure. As you will learn things will become much easier if the array is sorted.

A first try

Let's start by setting up a benchmark where we search through an unsorted array. The search procedure is of course a simple loop through all the elements in the array, let's call this method `search_unsorted`.

```
public static boolean search_unsorted(int[] array, int key) {  
    for (int index = 0; index < array.length ; index++) {  
        if (array[index] == key) {  
            return true;  
        }  
    }  
    return false;  
}
```

Set up the rest of a benchmark and do some measurements for a growing number of elements in the array. Describe the relationship between the size of the array and the time it takes to do the search. How long time does it take to search through an array of a million elements?

Now, if we know that the array is sorted we can of course do a quick optimization - we can stop the search once the next element in the array is larger then the key that we are looking for. Take a wild guess, how much better is this compared to our unsorted solution?

Set up a benchmark where you experiment with different data sizes. This is one way of generating a sorted array (no duplicates):

```
private static int[] sorted(int n) {
    Random rnd = new Random();
    int[] array = new int[n];
    int nxt = 0;
    for (int i = 0; i < n ; i++) {
        nxt += rnd.nextInt(10) + 1;
        array[i] = nxt;
    }
    return array;
}
```

How long time does it take to search through an array of a million entries? It might not sound very much but if our program constantly does search operations it will add up. There are however smarter things we can do and this will allow us to handle much larger data sets in reasonable time.

Binary search

The trick that we will do is what you would do if you searched for a word in a dictionary, a name in a phone book or a chapter in a book. You would not start on page one and look at one page at a time, you would jump to a page where you think it is likely to find the item and then jump a bit forward or backward depending on what you find on that page. After some jumps you have found the item you're looking for.

If we want to describe a general algorithm we could say that one should jump to the middle of the book and then jump either one quarter forward or backwards. In each round we will examine the page we jump to and make smaller and smaller jumps.

When we implement this algorithm we need to keep track off the *first possible page* and the *last possible page*. If we know this we can find the index in the middle and examine that page. If we find what we're looking for all is well but if not, we determine if we should update the *first* or *last* possible page. It could of course be the case that we have have no pages to jump to and in this case we know that the item that we are looking for is not in the book.

This algorithm is called *binary search* and is very efficient. There are of course many ways to encode it but here is some skeleton code that will get you started:

```
public static boolean binary_search(int[] array, int key) {

    int first = 0;
    int last = array.length-1;
```

```

while (true) {
    // jump to the middle
    int index = ..... ;

    if (array[index] == key) {
        // hmm what now?
        :
    }
    if (array[index] < key && index < last) {
        // The index position holds something that is less than
        // what we're looking for, what is the first possible page?
        first = ..... ;
        continue;
    }
    if (array[index] > key && index > first) {
        // The index position holds something that is larger than
        // what we're looking for, what is the last possible page?
        last = ..... ;
        continue;
    }
    // Why do we land here? What should we do?
    :
}
}

```

Re-run your benchmarks but now using the binary search. Report the execution time and describe a function that given the size of the array roughly describes the execution time. How long time does it take to search through an array of one million items? Without running an experiment - how long time do you estimate that it would take to search through an array of 64M items? Give it a try - how well did you estimate the execution time?

Even better

Remember the task you had in a previous assignment where you found all duplicates in two unsorted arrays. For each of the items in the first array you searched through the whole second array looking for duplicates. How would this strategy change if you would be given two sorted arrays?

If the arrays are sorted we could as before go through the first array, item by item, but we could now be smart when looking through the second array. We could use our implementation of the binary search algorithm and determine if a duplicate is found. Do some benchmark and present the improvement in run time that you get. Can we do even better?

Let's rewrite the algorithm and do as follows. Keep track of the *next element in the first array*. If the *next element in the second array* is smaller than the *next in the first*, then move forward in the second array. If it is equal (and then we found a duplicate) or greater then we move forward in the first array.

Assume that the two arrays themselves do not contain any duplicates (use the generator in the previous example). Run some benchmarks and compare the execution time with the run time using: unsorted arrays, binary search and your final version.

the cost of sorting

Working with sorted data is much more efficient than working with unsorted data. The problem is of course that not all data is sorted so the question is how hard it is to sort data. In the last experiment the question is if it is worth first sorting the two arrays before doing the search for duplicates. How expensive is it to sort an array and what algorithms do we have - this will be the question of the next assignment.